

# Multi-Port Memory with Bi-directional Ports for FPGAs Using XOR and LVT Methods

XXX  
XXX@XXX

## Abstract

We propose a generalization to a previous XOR memory implementation to allow for any number of full, read-only, and write-only ports. This implementation creates these addition ports by encoding and decoding the entries with XOR logic. Encoding the incoming (written) data with the data corresponding to the other ports and decoding the the outgoing read data by XORing all the entries in the memory banks. Using XOR operators to accomplish this encoding/decoding has advantages over other approaches because the XOR operator is a simple and efficient operation. The result is a high-throughput multi-ported memory that is particularly well-suited for implementation on FPGAs. This paper also presents an efficient architecture for creating live value table (LVT) memory using an XOR-based scheme by utilizing its bidirectional capabilities. We use an XOR memory with full ports to implement a bi-directional live value table design. We evaluate the architecture's performance and resource utilization, and show that the XOR-based bidirectional live value table is a compelling alternative for applications requiring high-performance, flexible memory access.

## ACM Reference Format:

XXX. 2025. Multi-Port Memory with Bi-directional Ports for FPGAs Using XOR and LVT Methods. In *Proceedings of International Symposium on Field Programmable Gate Arrays (FPGA '26)*. ACM, New York, NY, USA, 5 pages. <https://doi.org/XXXXXXX.XXXXXXX>

## 1 Motivation

As computation needs keep increasing, one way to keep up has been specialized architectures. FPGAs provide a way to implement architectures without taping out an ASIC. However, the limitations of FPGA resources requires some creativity to map designs to FPGAs. This paper explores how to overcome the limitation of FPGAs that have a limited number of ports. Specifically we propose a method to create memories with more than 2 ports.

The major FPGA vendors (AMD[13], Intel[5], Lattice[9], Microchip[10], and Achronix[2]) implement distributed memory (small memories) and block memory (large memories) differently. However they all share some characteristics. All vendors support distributed memory configurations with 1 full or write port and between 1 to 3 read ports. All vendors support block memory with 2 full ports. AMD has a limited offering of multiport memories, but otherwise none of the vendors support memories with more than 2

full ports. Although this limitation is problematic for designs requiring multiple ports (and particularly write or full ports), we show that these resources make it possible to achieve high throughput quad and octal full port memories.

## 2 Source Code

We provide all of the source code used in implementation and testing our design at <https://anonymous.4open.science/r/mpm-7666/>. We tested our design with Verilator[12] and synthesized the design with Vivado[3] targeting an AMD Virtex UltraScale+ HBM VU47P-3 FPGA.

## 3 Related Work

Several previous solutions to the port limit on FPGAs exist including multi-pumping, banking and replication[8]. Multi-pumping is the process of reducing the clock speed to increase the number of ports. For example, a 300Mhz single port memory can handle two 150Mhz ports. Banking[11] requires stalls and routing logic due to the segmented memory. Our design is most similar to replication. Replication involves tying the write ports of multiple memories together to create additional read ports.

## 4 XOR memory

We propose a simple generalization to XOR memories presented in previous work[6]. XOR memories work by using the following property:

$$a \oplus b \oplus b = a \quad (1)$$

Where  $\oplus$  is the bitwise XOR operator. Combined with the commutative and associative property of XOR we can generalize this property to:

$$b_0 \oplus \dots \oplus b_n (a \oplus b_0 \oplus \dots \oplus b_n) = a \quad (2)$$

We add bidirectional ports and analyze the performance of distributed memory and block memory versions of this design. We also present applications for these memories.

The number of RAMs (NRAM) needed is:

$$N_{RAM} = (W + F)(W + F + R) - W \quad (3)$$

Where  $W$  is the number of write ports,  $R$  is the number of read ports and  $F$  is the number of full ports. This expands to:

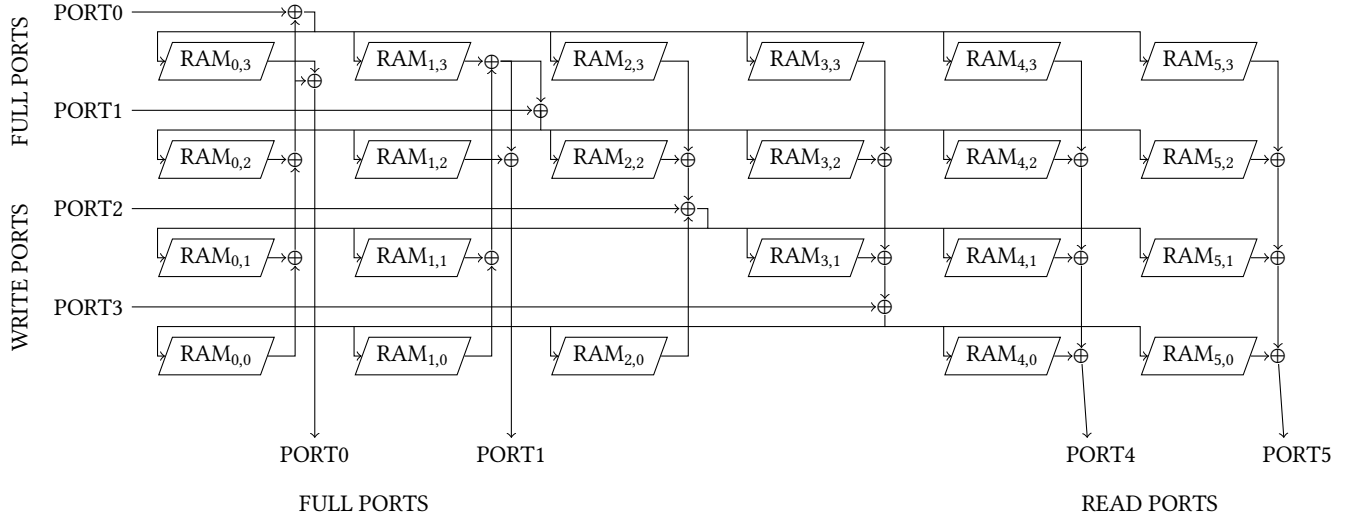
$$N_{RAM} = W^2 + 2WF + F^2 + WR + FR - W \quad (4)$$

Figure 1 demonstrates how to create an XOR memory with 2 read ports, 2 write ports and 2 full ports, which requires 22 RAMs. To read data from an XOR memory all of the data from the RAMs in one column is XORed from the same address. For example, say address  $x$  has values  $A, B, C$  and  $D$ , the data read would be  $A \oplus B \oplus C \oplus D$ .

Writing to the memory involves reading from all memories except the current row (say row/port 2 in figure 1) and XORing the

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).  
FPGA '26, Seaside, CA

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.  
ACM ISBN 978-1-4503-XXXX-X/2018/06  
<https://doi.org/XXXXXXX.XXXXXXX>



**Figure 1: A multi-port memory with 2 full ports, 2 write-only ports and 2 read-only ports.**

incoming data  $E$  (in the example this results in  $A \oplus B \oplus D \oplus E$ ) and storing that value in all the RAMs in that row.

The next time that data is read the result will be  $A \oplus B \oplus (A \oplus B \oplus D \oplus E) \oplus D$ , which equals  $E$ .

You may notice that writing to a port involves XORing all but one stored value and reading involves XORing all values. This enables full ports to be created just by adding one RAM to what would otherwise be just a write port.

Note that this memory requires that all of the RAMs in a row have the same data. Initializing the rams to the same data (e.g., all 0s) is required for the memory to operate properly. This is not an issue in FPGAs since the memory can be initialized to 0. since rows are written to at the same time, the same they will remain the same as long as the memories are initially.

We provide a second example with Figure 2 showing a clock cycle of an XOR memory with 2 full ports.

## 5 Analysis of XOR memory

We analyze several configurations of XOR memories. Particularly we vary the number of ports and width of the memory.

Table 1 we show the varying resources and frequency of the memory for different numbers of ports. As seen the number of LUTs used for memory is relatively large particularly for more ports.

As expected increasing the width of the memory resulted in a roughly linear increase in resource usage and a decrease in timing performance (see table 2).

Also as expected increasing the depth of the memory resulted in a roughly linear increase in resource usage and a decrease in timing performance (see table 3).

In table 4 we show an implementation using Block RAMs. This required pipelining and introducing a cycle of write delay. One could change the block ram to be write before read to remove the cycle of write delay.

**Table 1: Synthesis results of XOR memory for different port counts. The memory has a depth of 1024 and width of 32bits.**

Ports	LUTS	LUTS configured as memory	FF	BRAM	Max Frequency
2	2,706	2,304	0	0	625Mhz
4	11,556	9,728	0	0	500Mhz
8	48,328	39,936	0	0	357Mhz
16 <sup>1</sup>	196,640	161,792	0	0	XMhz
32 <sup>2</sup>	794,688	651,264	0	0	N/A

**Table 2: Synthesis results of XOR memory for different widths. The memory has a depth of 1024 and 8 ports.**

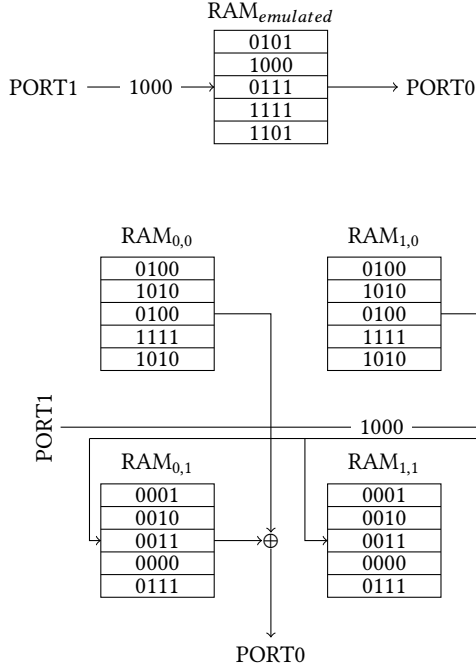
Width	LUTS	LUTS configured as memory	FF	BRAM	Max Frequency
1	2,032	1,920	0	0	526Mhz
2	4,040	3,840	0	0	500Mhz
4	8,797	7,680	0	0	400Mhz
8	12,140	9,984	0	0	385Mhz
16	24,201	19,968	0	0	370Mhz
32	48,328	39,936	0	0	357Mhz

## 6 Live Value Table Memory

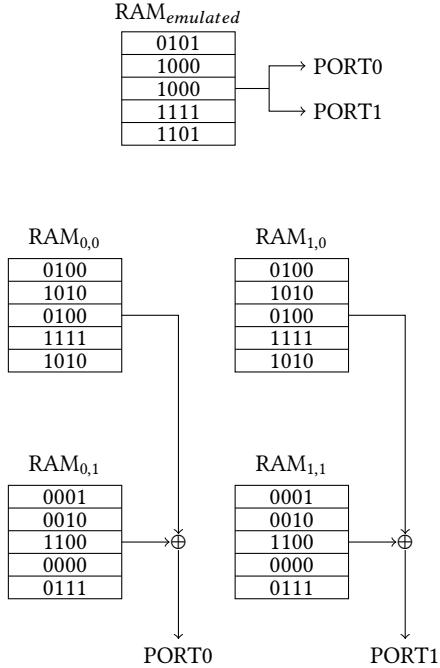
XOR memories can be used by themselves, however a live value table (LVT) may be more efficient.

We present a LVT memory that utilizes XOR memory.

Previous work used distributed memory [1]. However this work did not use bidirectional XOR ports in their implementation. In an alternate implementation[7], the live value table was implemented with registers.



(a) The memory on clock cycle 0, where  $1000_2$  is written to address 2 from port 1 and  $0111_2$  is read from address 2 from port 0.



(b) The memory on clock cycle 1, where  $1000_2$  is read from address 2 from both ports.

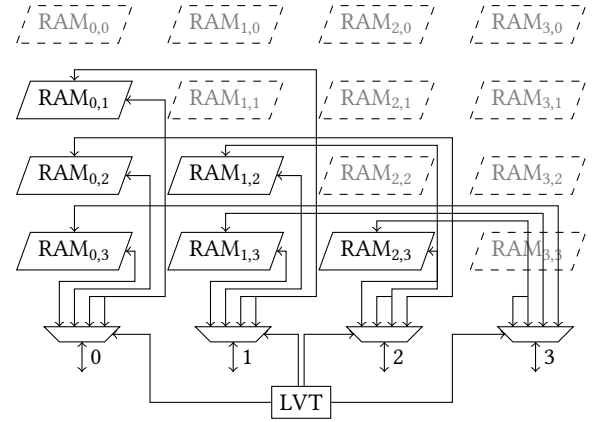
**Figure 2:** This shows an XOR memory during 2 clock cycles of operation. The XOR memory has 2 full ports and is constructed from 4 RAMs with 1 read and 1 write port.

**Table 3: Synthesis results of XOR memory for different depths. The memory has a width of 32 bits and 8 ports.**

Depth	LUTS	LUTS configured as memory	FF	BRAM	Max Frequency
32	2,016	1,248	0	0	588Mhz
64	3,264	2,496	0	0	556Mhz
128	6,483	4,992	0	0	526Mhz
256	12,565	9,984	0	0	476Mhz
512	24,621	19,968	0	0	455Mhz
1,024	48,328	39,936	0	0	357Mhz

**Table 4: Synthesis results of XOR memory with block RAMs for different port counts. The memory has a depth of 1024 and width of 32 bits.**

Ports	LUTS	LUTS configured as memory	FF	BRAM	Max Frequency
2	128	0	150	4	278Mhz
4	256	0	300	16	269Mhz
8	768	0	600	64	197Mhz
16 <sup>1</sup>	3,072	0	1,200	256	XMhz
32 <sup>1</sup>	10,240	0	2,400	1024	XMhz



**Figure 3:** Multi-port memory created with bi-directional dual-port memories. Note, a live value table is needed to determine which memory has the most recent value.

We create a LVT memory using the technique described in [4]. This live value memory is composed of 2-(full)port memories. Each port shares a RAM with another port. This results in  $F(F-1)/2$  RAMs being needed, where  $F$  is the number of bi-directional (full) ports. See figure 3.

The memory gets its name because of a multi-port memory that tracks the most recent stored value (aka a live value table). The point of a multi-port memory that requires a multi-port memory is that wide (e.g. 32 bit data) can be stored more efficiently this way.

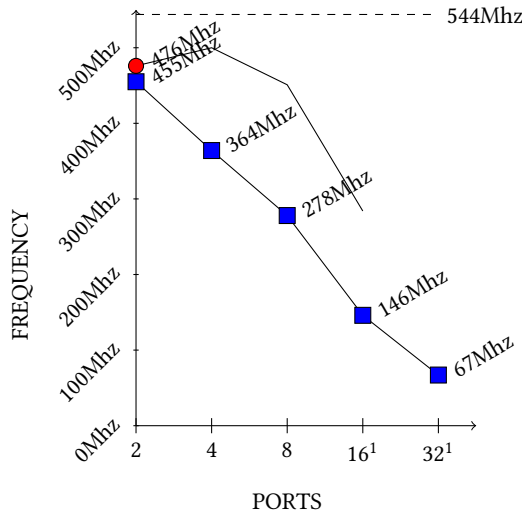


Figure 4: Frequency of LVT design.

Instead of using a register based live-value table as in [4] we use a xor memory similar to [1].

We show we utilize x% less resources than LVT and I-LVT.

## 7 Analysis of LVT Memory

We explore LVT designs with 2 to 32 ports. Although 16 and 32 port designs fit on large FPGAs, we believe smaller 4 and 8 port designs are more practical. We say more practical because of the high resource usage of XOR and LVT memories at high port counts.  $F^2$  for XOR and  $F(F-1)/2$  for LVT. However we were able to synthesize a 32 port memory. Higher port counts also had worse timing performance (see table 5 and figure 4).

To get better timing performance we created a pipelined version of the memory. This memory has major drawbacks: In addition to read delay the memory has write delay. The write delay means write conflicts occur on adjacent clock cycles not just current clock cycles. However we achieve better timing performance with this memory (see table 6 and figure 4).

Without write delay an 8 port memory runs at Xmhz (x% of max). With write delay and pipelining the design runs at Xmhz (x% of max).

## 8 Conclusion

XOR and LVT techniques can efficiently create memories with multiple ports. Although at the cost of using more memory space. Depending on the application, these memories may be the best option from the many available when creating a multi-port memory.

## References

- [1] Ameer M. S. Abdelhadi and Guy G. F. Lemieux. 2016. Modular Switched Multiported SRAM-Based Memories. *ACM Trans. Reconfigurable Technol. Syst.* 9, 3, Article 22 (July 2016), 26 pages. doi:10.1145/2851506
- [2] Achronix Semiconductor Corporation 2022. *Speedster7t FPGA Datasheet (DS015)*. Achronix Semiconductor Corporation. [https://www.achronix.com/sites/default/](https://www.achronix.com/sites/default/files/docs/Speedster7t_FPGA_Datasheet_DS015_8.pdf)

<sup>0</sup>To reduce the number of IO ports and fit the design on the FPGA we used a wrapper for the multi-port memory for designs with 16 and 32 ports.

Table 5: Synthesis results of LVT design for different port counts.

Ports	LUTS	LUTS configured as memory	FF	BRAM	Max Frequency
2	191	96	0	1	588Mhz
4	1,144	896	0	6	556Mhz
8	5,428	3,968	0	28	455Mhz
16 <sup>1</sup>	31,075	24,064	0	120	250Mhz
32 <sup>1</sup>	161,216	129,536	0	496	127Mhz

Table 6: Synthesis results of LVT design for different port counts for pipelined design

Ports	LUTS	LUTS configured as memory	FF	BRAM	Max Frequency
2	119	64	26	1	714Mhz
4	1,268	1,024	96	6	667Mhz
8	5,588	4,096	160	28	625Mhz
16 <sup>1</sup>	31,328	24,576	688	120	417Mhz
32 <sup>1</sup>	162,960	131,072	2,480	496	247Mhz

files/docs/Speedster7t\_FPGA\_Datasheet\_DS015\_8.pdf This document contains preliminary information and is subject to change without notice..

- [3] AMD. 2025. *Vivado Design Suite User Guide: Using the Vivado IDE* (2025.1 english ed.). AMD, San Jose, CA, USA. <https://docs.amd.com/r/en-US/ug893-vivado-ide>
- [4] Jongsok Choi, Kevin Nam, Andrew Canis, Jason Anderson, Stephen Brown, and Tomasz Czajkowski. 2012. Impact of Cache Architecture and Interface on Performance and Area of FPGA-Based Processor/Parallel-Accelerator Systems. In *2012 IEEE 20th International Symposium on Field-Programmable Custom Computing Machines*. 17–24. doi:10.1109/FCCM.2012.13
- [5] Intel Corporation 2024. *Intel Agilex 7 FPGA and SoC FPGA Datasheet*. Intel Corporation. <https://www.intel.com/content/www/us/en/programmable/support/literature/lit-agilex.html> Version 2024.09.03.
- [6] Charles Eric Laforest, Zimo Li, Tristan O'roure, Ming G. Liu, and J. Gregory Steffan. 2014. Composing Multi-Ported Memories on FPGAs. *ACM Trans. Reconfigurable Technol. Syst.* 7, 3, Article 16 (Sept. 2014), 23 pages. doi:10.1145/2629629
- [7] Charles Eric Laforest, Ming G. Liu, Emma Rae Rapati, and J. Gregory Steffan. 2012. Multi-ported memories for FPGAs via XOR. In *Proceedings of the ACM/SIGDA International Symposium on Field Programmable Gate Arrays* (Monterey, California, USA) (FPGA '12). Association for Computing Machinery, New York, NY, USA, 209–218. doi:10.1145/2145694.2145730
- [8] Charles Eric LaForest and J. Gregory Steffan. 2010. Efficient multi-ported memories for FPGAs. In *Proceedings of the 18th Annual ACM/SIGDA International Symposium on Field Programmable Gate Arrays* (Monterey, California, USA) (FPGA '10). Association for Computing Machinery, New York, NY, USA, 41–50. doi:10.1145/1723112.1723122
- [9] Lattice Semiconductor 2021. *FPGA-DS-02000 – Lattice iCE40 Family Data Sheet*. Lattice Semiconductor. [https://www.latticesemi.com/view\\_document?document\\_id=52424](https://www.latticesemi.com/view_document?document_id=52424) Accessed on 2025-09-08.
- [10] Microchip Technology Inc. 2025. *PolarFire and PolarFire SoC FPGA Fabric User Guide*. Microchip Technology Inc. [https://www.microchip.com/content/dam/mchp/documents/FPGA/ProductDocuments/UserGuides/PolarFire\\_PolarFire\\_SoC\\_FPGA\\_Fabric\\_User\\_Guide\\_VB.pdf](https://www.microchip.com/content/dam/mchp/documents/FPGA/ProductDocuments/UserGuides/PolarFire_PolarFire_SoC_FPGA_Fabric_User_Guide_VB.pdf) Document Number: UG0680.
- [11] Kevin R. Townsend, Osama G. Attia, Phillip H. Jones, and Joseph Zambreno. 2015. A Scalable Unsegmented Multiport Memory for FPGA-Based Systems. *International Journal of Reconfigurable Computing* 2015, 1 (2015), 826283. doi:10.1155/2015/826283 arXiv:https://onlinelibrary.wiley.com/doi/pdf/10.1155/2015/826283
- [12] Veripool, Inc. 2024. *Verilator Reference Manual*. Veripool, Inc. <https://verilator.org/guide/latest/> Version 5.041 (or latest as of your date of use). Available at <https://verilator.org/guide/latest/>.
- [13] Xilinx. 2020. *7 Series FPGAs Data Sheet: Overview*. Xilinx. <https://docs.amd.com/api/khuh/documents/2LByHkO~nSZXcei2D55fTg/content> v2.6.1.

Received 1 October 2025; revised XX XXX XXXX; accepted XX XXX  
XXXX